1 Announcements

Recommended Reading:

- CLRS Sec 2.2

- Fill out participation/reflection survey after SRE.

- Minus two Friday sections

- Plus Friday OH

- Plus Thursday evening section

- OH use: not for grading. Remember revision policy, and up to 4 R- grades on psets is still in A-range.

- Underloaded OH earlier in week, overloaded near deadlines.

Lecture Notes. We can post *drafts* of the detailed lecture notes on the course schedule in advance, but (a) strongly encourage you to work from the un-detailed ones during class in order to be actively engaged, and (b) note that the detailed lecture notes will usually be updated with corrections and more details after class.

2 Goals

So far, our conception of an algorithm has been informal: "a well-defined procedure for transforming inputs to outputs" whose runtime is measured as "basic operations" performed on a given input. This is unsatisfactory: how can we identify the fastest algorithm to solve a given problem if we don't have agreement on what counts as an algorithm or as a basic operation? To address this, we need to specify a *computational model* for describing algorithms.

What do we want from a computational model?

- Unambiguity. Precisely defines what is and isn't an algorithm, and what counts or doesn't count as a basic operation.

- Expressivity. Everything we intuitively consider to be an algorithm should be expressible, or we should have a good reason for excluding it.

- Mathematical simplicity. We want something that is manageable to reason about. In contrast, the specification of most modern programming languages is too long and cumbersome to prove anything general by hand.[1] For example, the Python Language Reference (Release 3.4.3), which "describes the syntax and 'core semantics' of the language" (not including libraries) is 135 pages long.

- Robustness. Small changes to the model should not fundamentally change what can or cannot be computed, and should have only a small effect on efficiency.

- Technological relevance. The model should capture reasonably well the implementability and efficiency of algorithms on actual computing hardware, even as technology evolves with time.

3   The RAM Model

Our first attempt at a precise model of computation is the *RAM model*, which models memory as an infinite array M of *natural numbers*.

Definition 3.1 (RAM Programs: syntax)**.** A *RAM Program* $P = (V, C_0, \ldots, C_{\ell-1})$ consists of a finite set V of *variables* (or *registers*), and a sequence $C_0, C_1, \ldots, C_{\ell-1}$ of *commands* (or *lines of code*), each chosen from the following:

- (assignment to a constant) `var` $= c$, for a variable `var` $\in V$ and a constant $c \in \mathbb{N}$.

- (arithmetic) $\texttt{var}_0 = \texttt{var}_1$ `op` $\texttt{var}_2$, for variables $\texttt{var}_0, \texttt{var}_1, \texttt{var}_2 \in V$, and an operation `op` chosen from $+, -, \times, /$.

- (read from memory) $\texttt{var}_0 = M[\texttt{var}_1]$ for variables $\texttt{var}_0, \texttt{var}_1 \in V$.

- (write to memory) $M[\texttt{var}_0] = \texttt{var}_1$ for variables $\texttt{var}_0, \texttt{var}_1 \in V$.

- (conditional goto) IF `var` $== 0$, GOTO k, where $k \in \{0, 1, \ldots, \ell\}$.

In addition, we require that every RAM Program has three special variables: `input_len`, `output_ptr`, and `output_len`.

Definition 3.2 (Computation of a RAM Program: semantics)**.** A RAM Program $P = (V, (C_0, \ldots, C_{\ell-1}))$ *computes* on an input x is as follows:

1. Initialization: The input x is encoded (in some predefined manner) as a sequence of natural numbers placed into memory locations $(M[0], \ldots, M[n-1])$, and all of the remaining memory locations are set to 0. The variable `input_len` is initialized to n, the length of x's encoding. All other variables are initialized to 0.

2. Execution: The sequence of commands $C_0, C_1, C_2, \ldots$ are executed in order (except when jumps are done due to GOTO commands), updating the values of variable and memory locations according to the usual interpretations of the operations. Since we are working with natural numbers, if the result of subtraction would be negative, it is replaced with 0. Similarly, the results of division are rounded down, and divide by 0 results in 0.

---

[1] As we'll discuss at the end of the course, there are software tools that assist in proving properties of programs. However, to make effective use of these tools you need to have experience in proving things about programs yourself!

3. Output: If line ℓ is reached (possibly due to a GOTO ℓ), the output $P(x)$ is defined to be the subarray of M of length output_len starting at location output_ptr. That is,

$$P(x) = (M[\texttt{output\_ptr}], M[\texttt{output\_ptr}+1], \ldots, M[\texttt{output\_ptr} + \texttt{output\_len} - 1]).$$

The *running time* of P on input x, denoted $\text{Time}_P(x)$, is defined to be: the number of commands executed during the computation (possibly ∞).

The definition of the RAM Model above is mathematically precise, so achieves our unambiguity desideratum (unless we've forgotten to specify something!).

The RAM Model also does quite well on the mathematical simplicity front. We described it in one page of lecture notes, compared to 100+ pages for most modern programming languages. That said, there are even simpler models of computation, such as the Turing Machine and the Lambda Calculus. However, those are harder to describe algorithms in and less accurately describe computing technology. We will briefly discuss those later in the course when we cover the Church–Turing Thesis, and they (along with other models of computation) are studied in depth in CS121.

Our focus for the rest of today's class will be to get convinced of the *expressivity* of the RAM model. We will do this by seeing how to implement algorithms we have seen in the RAM model. We will turn to robustness and technological relevance next time.

4   Iterative Algorithms

Let's see an example RAM program for Insertion Sort, when the keys and values are both given as natural numbers.

Input	: An array $x = (K_0, V_0, K_1, V_1, \ldots, K_{n-1}, V_{n-1})$, occupying memory locations $M[0], \ldots, M[2n-1]$	
Output	: A valid sorting of x. in the same memory locations as the input	
Variables:	`input_len, output_len, zero, one, two, output_ptr, outer_key_ptr, outer_rem,` `outer_key, inner_key_ptr, inner_rem, inner_key, key_diff, insert_key,` `insert_value, temp_ptr, temp_key, temp_value`	

0 `zero` $= 0$;   /* useful constants */

1 `one` $= 1$;

2 `two` $= 2$;

3 `output_ptr` $= 0$;   /* output will overwrite input */

4 `output_len` $=$ `input_len` $+$ `zero`;

5 `outer_key_ptr` $= 0$;   /* pointer to the key we want to insert */

6 `outer_rem` $=$ `input_len`$/$`two` ;   /* # outer-loop iterations remaining */

7   `outer_key_ptr` $=$ `outer_key_ptr` $+$ `two` ;   /* start of outer loop */

8   `outer_rem` $=$ `outer_rem` $-$ `one`;

9   IF `outer_rem` $== 0$ GOTO 34;

10   `outer_key` $= M[$`outer_key_ptr`$]$;   /* key to be inserted */

11   `inner_key_ptr` $= 0$;   /* pointer to potential insertion point */

12   `inner_rem` $=$ `outer_key_ptr`$/$`two` ;   /* # inner-loop iterations remaining */

13   `inner_key` $= M[$`inner_key_ptr`$]$;   /* start 1st inner loop */

14   `key_diff` $=$ `inner_key` $-$ `outer_key` ;   /* if inner_key $\leq$ outer_key, then */

15   IF `key_diff` $== 0$ GOTO 30 ;   /* proceed to next inner iteration */

16   `insert_key` $=$ `outer_key` $+$ `zero` ;   /* key to be inserted */

17   `temp_ptr` $=$ `outer_key_ptr` $+$ `one`;

18   `insert_value` $= M[$`temp_ptr`$]$;   /* value to be inserted */

19   `temp_key` $= M[$`inner_key_ptr`$]$;   /* start of 2nd inner loop */

20   `temp_ptr` $=$ `inner_key_ptr` $+$ `one`;

21   `temp_value` $= M[$`temp_ptr`$]$;

22 $M[$`inner_key_ptr`$]$ $=$ `insert_key`;

23 $M[$`temp_ptr`$]$ $=$ `insert_value`;

24   `insert_key` $=$ `temp_key` $+$ `zero`;

25   `insert_value` $=$ `temp_value` $+$ `zero`;

26   `inner_key_ptr` $=$ `inner_key_ptr` $+$ `two`;

27   IF `inner_rem` $== 0$ GOTO 7;

28   `inner_rem` $=$ `inner_rem` $-$ `one`;

29   IF `zero` $== 0$ GOTO 19;

30   `inner_key_ptr` $=$ `inner_key_ptr` $+$ `two`;

31   `inner_rem` $=$ `inner_rem` $-$ `one`;

32   IF `inner_rem` $== 0$ GOTO 7;

33   IF `zero` $== 0$ GOTO 13;

34 HALT ;   /* not an actual command */

**Algorithm 1:** RAM implementation of Insertion Sort

5   Data

Implicit in the expressivity requirement is that we can describe the inputs and outputs of algorithms in the model. In the RAM model, all inputs and outputs are arrays of natural numbers. How can we represent other types of data?

- (Signed) integers: natural numbers with an additional sign bit

- Rational numbers: pair of integers (numerator and denominator)

- Real numbers: impossible! (For those of you who have taken CS20, this is because the set of real numbers is uncountable, but the set of finite sequences of natural numbers is countable.) Although this is a constraint on expressivity, the inability to represent and manipulate infinite-precision real numbers is generally accepted as part of both the intuitive and physically realizable concepts of computation.[2]

  To approximate real numbers, Python and other programming languages often use *floating-point numbers*, in the form $x \cdot 2^y$ for signed integers x and y of bounded bit-length. Almost every operation on floating-point numbers introduces error through rounding, and it can be very tricky to reason about how these errors propagate. Thus, *in the rest of the course, we will generally stick to problems involving integer or rational arithmetic.* Dealing with floating-point errors is a topic sometimes covered in courses on numerical analysis, scientific computation, numerical linear algebra, and/or optimization, which you might find in the Applied Math or Applied Computation parts of the catalogue.

- Strings: an array of ASCII values (where each character is represented as a number in $\{0, 1 \ldots, 127\}$)

What about a fancy data structure like a binary search tree? We can represent a BST as an array of 4-tuples (K_0, V_0, P_L, R_R) where P_L and P_R are pointers to the left and right children. Let's consider the example from last class:

---

[2]That said, one can define real-number models of computation, and these are mathematically interesting to study. See the book *Complexity and Real Computation* by Blum, Shub, and Smale.

Assuming all of the associated values are 0, this would be represented as the following array of length 28:

$$[45, 0, 4, 8, 29, 0, 12, 16, 71, 0, 20, 0, 20, 0, 0, 0, 41, 0, 0, 0, 50, 0, 0, 24, 63, 0, 0, 0]$$

For nodes that do not have a left or right child, we assign the value of 0 to P_L or P_R. Assigning the pointer value to 0 does not refer to the value at the memory location of 0, as we assume that the root of the tree cannot be a child for any of the nodes in the rest of the tree. Note that there are many ways to construct a binary tree using this array representation.

6 Recursive Algorithms

We will not cover this in lecture, but include it here in case you are interested and/or want more convincing about the expressivity of the RAM Model.

It is not entirely obvious that the RAM Model can implement recursion, since there are no function calls in its description. The way this is done (both in theory and in practice) via the use of a *stack* data structure. We simulate a function call $f(x)$ through the following steps:

1. Push local variables (in scope of the calling code), the input x, and an indicator of which line number to return to after the f is done executing.

2. GOTO the line number that starts the implementation of f.

3. The implementation of f pops its input x off the top of the stack, computes $y = f(x)$, and pushes y onto the top of the stack, and then GOTO to the line number after the function call (which it also read off the top of the stack).

4. After the return, $y = f(x)$ is read off the top of the stack, along with the local variables needed to continue the computation where it left off before calling f.

Below we present an example for a recursive computation of the height of a binary tree. Since our RAM model doesn't allow negative numbers, our recursive functions will compute height plus one (so that an empty tree has height 0), and we will subtract one from the height at the end. Also, because this algorithm does not use any memory other than the stack and the arrays, we implement the stack as a contiguous segment of memory starting after the input. However, in general, one may need to implement it as a linked list in order to be able to skip over portions of memory that

are being used for global state.

Input   : A Binary Tree of Key-Value Pairs, given as an array of 4-tuples (K, V, P_L, P_R)
Output   : The height of the input tree

```
 0  zero = 0 ;                                          /* useful constants */
 1  one = 1;
 2  two = 2;
 3  stack_ptr = input_len + zero;
 4  M[stack_ptr] = zero ;         /* push pointer to root of tree to top of stack */
 5  stack_ptr = stack_ptr + one;
 6  M[stack_ptr] = zero ;                     /* branch-indicator for root call */
 7      branch = M[stack_ptr] ;                        /* pop branch indicator */
 8      stack_ptr = stack_ptr − one;
 9      node_ptr = M[stack_ptr] ;               /* pop pointer to current node */
10      stack_ptr = stack_ptr + two ;      /* and repush both back onto stack */
11      temp_ptr = node_ptr + two;
12      child_ptr = M[temp_ptr] ;                     /* pointer to left child */
13      IF child_ptr == 0 GOTO 22;
14      M[stack_ptr] = child_ptr ;               /* push pointer to left child */
15      stack_ptr = stack_ptr + one;
16      M[stack_ptr] = one ;                        /* left branch indicator */
17      IF zero == 0 GOTO 7 ;                                    /* recurse */
```

Algorithm 2: RAM implementation of Calculate Height

```
18     left_height = M[stack_ptr] ;                      /* pop height of left child */
19     stack_ptr = stack_ptr − one;
20     node_ptr = M[stack_ptr] ;                      /* pop pointer to current node */
21     IF zero == 0 GOTO 23

22     left_height = 0 ;                                      /* left child is empty */

23     temp_ptr = node_ptr + three;
24     child_ptr = M[temp_ptr] ;                          /* pointer to right child */
25     IF child_ptr == 0 GOTO 37;
26     M[stack_ptr] = left_height ;                    /* push height of left child */
27     stack_ptr = stack_ptr + one;
28     M[stack_ptr] = child_ptr ;                     /* push pointer to right child */
29     stack_ptr = stack_ptr + one;
30     M[stack_ptr] = two ;                              /* right branch indicator */
31     IF zero == 0 GOTO 7 ;                                           /* recurse */

32     right_height = M[stack_ptr] ;                  /* pop height of right child */
33     stack_ptr = stack_ptr − one;
34     left_height = M[stack_ptr] ;                    /* pop height of left child */
35     stack_ptr = stack_ptr − one;
36     IF zero == 0 GOTO 38;

37     right_height = 0
38     branch = M[stack_ptr] ;                             /* pop branch indicator */
39     diff_heights = left_height − right_height;
40     IF diff_heights == 0 GOTO 43 ;                      /* right-child taller */
41     height = left_height + one ;                         /* left-child taller */
42     IF zero == 0 GOTO 44;
43     height = right_height + one;

44     IF branch == zero GOTO 50;
45     M[stack_ptr] = height ;                              /* push return value */
46     branch = branch − one;
47     IF branch == zero GOTO 18;
48     branch = branch − one;
49     IF branch == zero GOTO 32;
50 height = height − one ;                        /* subtract one for output height */
51 M[stack_ptr] = height;
52 output_ptr = stack_ptr;
53 output_len = 1;
54 HALT
```

Algorithm 3: RAM implementation of Calculate Height (cont.)

7   Reductions

We may not have time to cover this in class, but is again included for your interest.

We can also formalize *reductions* using the following extension of the RAM model.

Definition 7.1. An *oracle-aided RAM Program* is like an ordinary RAM program, except it can also have commands of the form

$$\text{ORACLE}(\text{var}_0, \text{var}_1, \text{var}_2),$$

which means call the oracle on the array $(M[\text{var}_0], M[\text{var}_0 + 1], \ldots, M[\text{var}_0 + \text{var}_1 - 1])$ and write the oracle's answer in the locations $(M[\text{var}_2], M[\text{var}_2 + 1], \ldots)$.

For example, our reduction from IntervalScheduling-Decision to Sorting is given by the following oracle-aided RAM program:

```
Input     : An array x = (a_0, b_0, a_1, b_1, ..., a_{n-1}, b_{n-1}), occupying memory locations
            M[0], ..., M[2n - 1], with a_i ≤ b_i for all i
Output    : 1 (YES) if all of the intervals [a_i, b_i] are disjoint, 0 (NO) otherwise
0 zero = 0;
1 one = 1;
2 two = 2;
3 ORACLE(zero, input_len, zero) ;                    /* sort input by start time */
4 output_ptr = 0;
5 output_len = 1;
6 M[zero] = one ;                                    /* default output is 1 = YES */
7 temp_ptr = 1;
8 remaining = input_len − two ;  /* how many adjacent pairs left to check, times
   two */

9     IF remaining == 0 GOTO 19;
10    end_curr = M[temp_ptr] ;              /* read end time of current interval */
11    temp_ptr = temp_ptr + one;
12    start_next = M[temp_ptr] ;            /* read start time of next interval */
13    temp = start_next − end_curr;
14    IF temp == 0 GOTO 18 ;                             /* conflict found */
15    temp_ptr = temp_ptr + one;
16    remaining = remaining − two;
17    IF zero == 0 GOTO 9;

18 M[zero] = zero ;                                  /* change output to 0 = NO */
19 HALT
```

Algorithm 4: Oracle-RAM implementation of IntervalScheduling-Decision$\leq_{O(n),n}$ Sorting